

# CompVis\_Assignment2\_MarcoFuchs

January 14, 2022

## 1 MSE Computer Vision - Assignment 2

### 1.1 Read image paths

```
[1]: import os

TRAIN_PATH = 'data/train2/'
VAL_PATH = 'data/val/'
TST_PATH = 'data/test/'

def path_names_list(path, name_snipped):
    return sorted([os.path.join(path, fname)
                    for fname in os.listdir(path)
                    if fname.endswith(".png") and name_snipped in fname])

train_img_paths = path_names_list(TRAIN_PATH, "train_img")
train_lbl_paths = path_names_list(TRAIN_PATH, "train_lbl")
val_img_paths = path_names_list(VAL_PATH, "val_img")
val_lbl_paths = path_names_list(VAL_PATH, "val_lbl")
tst_img_paths = path_names_list(TST_PATH, "test_img")
tst_lbl_paths = path_names_list(TST_PATH, "test_lbl")

print("Length training data:", len(train_img_paths), len(train_lbl_paths))
for n, (img_path, lbl_path) in enumerate(zip(train_img_paths, train_lbl_paths)):
    print(f"{img_path} -> {lbl_path}")
    if n == 3:
        print()
        break

print("Length validation data:", len(val_img_paths), len(val_lbl_paths))
for n, (img_path, lbl_path) in enumerate(zip(val_img_paths, val_lbl_paths)):
    print(f"{img_path} -> {lbl_path}")
    if n == 3:
        print()
        break

print("Length test data:", len(tst_img_paths), len(tst_lbl_paths))
```

```

for n, (img_path, lbl_path) in enumerate(zip(tst_img_paths, tst_lbl_paths)):
    print(f"{img_path} -> {lbl_path}")
    if n == 3:
        break

```

Length training data: 23520 23520

```

data/train2/train_img_000_00.png -> data/train2/train_lbl_000_00.png
data/train2/train_img_000_01.png -> data/train2/train_lbl_000_01.png
data/train2/train_img_000_02.png -> data/train2/train_lbl_000_02.png
data/train2/train_img_000_03.png -> data/train2/train_lbl_000_03.png

```

Length validation data: 8544 8544

```

data/val/val_img_000_00.png -> data/val/val_lbl_000_00.png
data/val/val_img_000_01.png -> data/val/val_lbl_000_01.png
data/val/val_img_000_02.png -> data/val/val_lbl_000_02.png
data/val/val_img_000_03.png -> data/val/val_lbl_000_03.png

```

Length test data: 1888 1888

```

data/test/test_img_000_00.png -> data/test/test_lbl_000_00.png
data/test/test_img_000_01.png -> data/test/test_lbl_000_01.png
data/test/test_img_000_02.png -> data/test/test_lbl_000_02.png
data/test/test_img_000_03.png -> data/test/test_lbl_000_03.png

```

## 1.2 Dataloader

```

[27]: import numpy as np
import tensorflow as tf
from sklearn.utils import shuffle

IMG_SHAPE = 128

def normalize(image, mask):
    image = tf.cast(image, tf.float32) / 255.0
    mask += 1 # to avoid negative labels
    return image, mask

def load_and_preprocess(img_filepath, mask_filepath):
    img = tf.io.read_file(img_filepath)
    img = tf.io.decode_jpeg(img, channels=3)
    img = tf.image.resize(img, [IMG_SHAPE, IMG_SHAPE])

    mask = tf.io.read_file(mask_filepath)
    mask = tf.io.decode_png(mask, channels=1)
    mask = tf.image.resize(mask, [IMG_SHAPE, IMG_SHAPE])

    img, mask = normalize(img, mask)

    return img, mask

```

```

AUTO = tf.data.experimental.AUTOTUNE
BATCH_SIZE = 32

# prepare data loaders
train_img_paths, train_lbl_paths = shuffle(train_img_paths, train_lbl_paths,
↳random_state=42)
val_img_paths, val_lbl_paths = shuffle(val_img_paths, val_lbl_paths,
↳random_state=42)
tst_img_paths, tst_lbl_paths = shuffle(tst_img_paths, tst_lbl_paths,
↳random_state=42)

trainloader = tf.data.Dataset.from_tensor_slices((train_img_paths,
↳train_lbl_paths))
valloader = tf.data.Dataset.from_tensor_slices((val_img_paths, val_lbl_paths))
tstloader = tf.data.Dataset.from_tensor_slices((tst_img_paths, tst_lbl_paths))

trainloader = (
    trainloader
    .shuffle(1024)
    .map(load_and_preprocess, num_parallel_calls=AUTO)
    .batch(BATCH_SIZE)
    .prefetch(AUTO) # prefetch next batch in a thread
)

valloader = (
    valloader
    .map(load_and_preprocess, num_parallel_calls=AUTO)
    .batch(BATCH_SIZE)
    .prefetch(AUTO)
)

tstloader = (
    tstloader
    .map(load_and_preprocess, num_parallel_calls=AUTO)
    .batch(BATCH_SIZE)
    .prefetch(AUTO)
)

```

```

[3]: labels = [
    # name                                id  trainId  category                catId  ↳
    ↳hasInstances  ignoreInEval  color
    [ 'unlabeled' , 0 ,      255 , 'void' , 0 ,
    ↳False , True , ( 0, 0, 0) ],

```

|                          |                     |                      |     |        |
|--------------------------|---------------------|----------------------|-----|--------|
| [ 'ego vehicle'          | , 1 ,               | 255 , 'void'         | , 0 | ,<br>␣ |
| ↪False , True            | , ( 0, 0, 0 ) ],    |                      |     |        |
| [ 'rectification border' | , 2 ,               | 255 , 'void'         | , 0 | ,<br>␣ |
| ↪False , True            | , ( 0, 0, 0 ) ],    |                      |     |        |
| [ 'out of roi'           | , 3 ,               | 255 , 'void'         | , 0 | ,<br>␣ |
| ↪False , True            | , ( 0, 0, 0 ) ],    |                      |     |        |
| [ 'static'               | , 4 ,               | 255 , 'void'         | , 0 | ,<br>␣ |
| ↪False , True            | , ( 0, 0, 0 ) ],    |                      |     |        |
| [ 'dynamic'              | , 5 ,               | 255 , 'void'         | , 0 | ,<br>␣ |
| ↪False , True            | , (111, 74, 0 ) ],  |                      |     |        |
| [ 'ground'               | , 6 ,               | 255 , 'void'         | , 0 | ,<br>␣ |
| ↪False , True            | , ( 81, 0, 81 ) ],  |                      |     |        |
| [ 'road'                 | , 7 ,               | 0 , 'flat'           | , 1 | ,<br>␣ |
| ↪False , False           | , (128, 64,128) ],  |                      |     |        |
| [ 'sidewalk'             | , 8 ,               | 1 , 'flat'           | , 1 | ,<br>␣ |
| ↪False , False           | , (244, 35,232) ],  |                      |     |        |
| [ 'parking'              | , 9 ,               | 255 , 'flat'         | , 1 | ,<br>␣ |
| ↪False , True            | , (250,170,160) ],  |                      |     |        |
| [ 'rail track'           | , 10 ,              | 255 , 'flat'         | , 1 | ,<br>␣ |
| ↪False , True            | , (230,150,140) ],  |                      |     |        |
| [ 'building'             | , 11 ,              | 2 , 'construction'   | , 2 | ,<br>␣ |
| ↪False , False           | , ( 70, 70, 70 ) ], |                      |     |        |
| [ 'wall'                 | , 12 ,              | 3 , 'construction'   | , 2 | ,<br>␣ |
| ↪False , False           | , (102,102,156) ],  |                      |     |        |
| [ 'fence'                | , 13 ,              | 4 , 'construction'   | , 2 | ,<br>␣ |
| ↪False , False           | , (190,153,153) ],  |                      |     |        |
| [ 'guard rail'           | , 14 ,              | 255 , 'construction' | , 2 | ,<br>␣ |
| ↪False , True            | , (180,165,180) ],  |                      |     |        |
| [ 'bridge'               | , 15 ,              | 255 , 'construction' | , 2 | ,<br>␣ |
| ↪False , True            | , (150,100,100) ],  |                      |     |        |
| [ 'tunnel'               | , 16 ,              | 255 , 'construction' | , 2 | ,<br>␣ |
| ↪False , True            | , (150,120, 90) ],  |                      |     |        |
| [ 'pole'                 | , 17 ,              | 5 , 'object'         | , 3 | ,<br>␣ |
| ↪False , False           | , (153,153,153) ],  |                      |     |        |
| [ 'polegroup'            | , 18 ,              | 255 , 'object'       | , 3 | ,<br>␣ |
| ↪False , True            | , (153,153,153) ],  |                      |     |        |
| [ 'traffic light'        | , 19 ,              | 6 , 'object'         | , 3 | ,<br>␣ |
| ↪False , False           | , (250,170, 30) ],  |                      |     |        |
| [ 'traffic sign'         | , 20 ,              | 7 , 'object'         | , 3 | ,<br>␣ |
| ↪False , False           | , (220,220, 0) ],   |                      |     |        |
| [ 'vegetation'           | , 21 ,              | 8 , 'nature'         | , 4 | ,<br>␣ |
| ↪False , False           | , (107,142, 35) ],  |                      |     |        |
| [ 'terrain'              | , 22 ,              | 9 , 'nature'         | , 4 | ,<br>␣ |
| ↪False , False           | , (152,251,152) ],  |                      |     |        |

```

    [ 'sky' , 23 , 10 , 'sky' , 5 ,
    ↪False , False , ( 70,130,180) ],
    [ 'person' , 24 , 11 , 'human' , 6 ,
    ↪True , False , (220, 20, 60) ],
    [ 'rider' , 25 , 12 , 'human' , 6 ,
    ↪True , False , (255, 0, 0) ],
    [ 'car' , 26 , 13 , 'vehicle' , 7 ,
    ↪True , False , ( 0, 0,142) ],
    [ 'truck' , 27 , 14 , 'vehicle' , 7 ,
    ↪True , False , ( 0, 0, 70) ],
    [ 'bus' , 28 , 15 , 'vehicle' , 7 ,
    ↪True , False , ( 0, 60,100) ],
    [ 'caravan' , 29 , 255 , 'vehicle' , 7 ,
    ↪True , True , ( 0, 0, 90) ],
    [ 'trailer' , 30 , 255 , 'vehicle' , 7 ,
    ↪True , True , ( 0, 0,110) ],
    [ 'train' , 31 , 16 , 'vehicle' , 7 ,
    ↪True , False , ( 0, 80,100) ],
    [ 'motorcycle' , 32 , 17 , 'vehicle' , 7 ,
    ↪True , False , ( 0, 0,230) ],
    [ 'bicycle' , 33 , 18 , 'vehicle' , 7 ,
    ↪True , False , (119, 11, 32) ],
    [ 'license plate' , -1 , -1 , 'vehicle' , 7 ,
    ↪False , True , ( 0, 0,142) ],
]

# segmentation_classes = [label[0] for label in labels]
SEGMENTATION_CLASSES_DICT = {n: label[0] for n, label in enumerate(labels) if
    ↪label[0] != "license plate"}
SEGMENTATION_CLASSES_DICT[-1] = "license plate"

OUTPUT_CHANNEL = len(SEGMENTATION_CLASSES_DICT)
print(f"There are {OUTPUT_CHANNEL} segmentatin classes.")
print(SEGMENTATION_CLASSES_DICT)

```

There are 35 segmentatin classes.

```

{0: 'unlabeled', 1: 'ego vehicle', 2: 'rectification border', 3: 'out of roi',
4: 'static', 5: 'dynamic', 6: 'ground', 7: 'road', 8: 'sidewalk', 9: 'parking',
10: 'rail track', 11: 'building', 12: 'wall', 13: 'fence', 14: 'guard rail', 15:
'bridge', 16: 'tunnel', 17: 'pole', 18: 'polegroup', 19: 'traffic light', 20:
'traffic sign', 21: 'vegetation', 22: 'terrain', 23: 'sky', 24: 'person', 25:
'rider', 26: 'car', 27: 'truck', 28: 'bus', 29: 'caravan', 30: 'trailer', 31:
'train', 32: 'motorcycle', 33: 'bicycle', -1: 'license plate'}

```


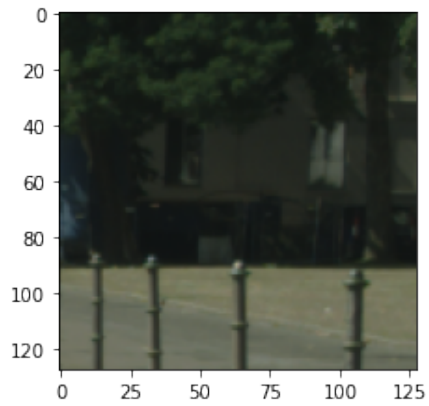



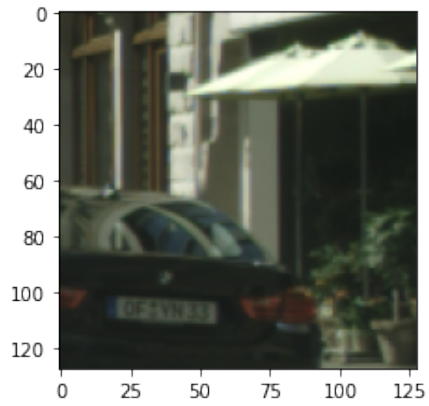
### 1.3 Visualize

```
[4]: import matplotlib.pyplot as plt

num_example_to_display = 3
img, mask = next(iter(valloader))
fig, axs = plt.subplots(nrows=num_example_to_display, ncols=2,
    figsize=(12,4*num_example_to_display))

for i in range(num_example_to_display):
    axs[i][0].imshow(img[i])
    axs[i][1].imshow(np.squeeze(mask[i],-1), cmap='gray');
```

2022-01-13 09:42:10.813714: I  
tensorflow/compiler/mlir/mlir\_graph\_optimization\_pass.cc:185] None of the MLIR  
Optimization Passes are enabled (registered 2)



## 1.4 Model

```
[5]: from tensorflow.keras.layers import *
     from tensorflow.keras.models import *
     import keras
```

```
[6]: def fcn_simple_no_border(input_height:int, input_width:int) -> keras.Model:
     """
```

```

Create a simple fcn model for semantic segmentation with 2 classes
"""
model = keras.Sequential()

# we use grayscale (1-channel input)

# (used to define input shape on the first layers)
model.add(keras.layers.Layer(input_shape=(input_height , input_width, 3)))

# add 3 convolutional layers with 3x3 filters
model.add(keras.layers.Convolution2D(filters=4, kernel_size=3, strides=(2,↵
↵2), padding='same', activation='relu'))
model.add(keras.layers.Conv2DTranspose(filters=4, kernel_size=3,↵
↵strides=(2, 2), padding='same', activation='relu'))
model.add(keras.layers.Convolution2D(filters=4, kernel_size=3,↵
↵padding='same', activation='relu'))

# go to logits which is the number of classes and add sigmoid layer for↵
↵activation
model.add(keras.layers.Convolution2D(filters=1, kernel_size=1,↵
↵activation=None,
                                kernel_initializer=keras.initializers.
↵TruncatedNormal(mean=0.0, stddev=0.001, seed=None)))
model.add(keras.layers.Activation('sigmoid'))

# reshape so that we have a sample for each pixel
model.add(keras.layers.Reshape(target_shape=(input_height, input_width, 1)))

return model

```

```

[7]: def model_u(img_shape, output_channels):
    model = keras.Sequential()
    model.add(keras.layers.Layer(input_shape=(img_shape , img_shape, 3)))

    # down
    model.add(Conv2D(filters=32, kernel_size=(3, 3), activation='relu',↵
↵padding='same'))
    model.add(Conv2D(filters=32, kernel_size=(3, 3), activation='relu',↵
↵padding='same'))
    model.add(MaxPooling2D((2, 2)))

    model.add(Conv2D(filters=64, kernel_size=(3, 3), activation='relu',↵
↵padding='same'))
    model.add(Conv2D(filters=64, kernel_size=(3, 3), activation='relu',↵
↵padding='same'))
    model.add(MaxPooling2D((2, 2)))

```



```

        model.add(Conv2D(filters=128, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
        model.add(Conv2D(filters=128, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
        model.add(MaxPooling2D((2, 2)))

#         model.add(Conv2D(filters=256, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
#         model.add(Conv2D(filters=256, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
#         model.add(MaxPooling2D((2, 2)))

        # up
#         model.add(Conv2D(filters=512, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
#         model.add(Conv2D(filters=512, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
#         model.add(Conv2DTranspose(filters=256, kernel_size=(2,2), strides=(2,2),
        ↪padding='same'))

        model.add(Conv2D(filters=256, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
        model.add(Conv2D(filters=256, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
        model.add(Conv2DTranspose(filters=128, kernel_size=(2,2), strides=(2,2),
        ↪padding='same'))

        model.add(Conv2D(filters=128, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
        model.add(Conv2D(filters=128, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
        model.add(Conv2DTranspose(filters=64, kernel_size=(2,2), strides=(2,2),
        ↪padding='same'))

        model.add(Conv2D(filters=64, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
        model.add(Conv2D(filters=64, kernel_size=(3, 3), activation='relu',
        ↪padding='same'))
        model.add(Conv2DTranspose(filters=32, kernel_size=(2,2), strides=(2,2),
        ↪padding='same'))

        model.add(Conv2D(output_channels, (3, 3), activation='softmax',
        ↪padding='same'))

    return model

```

## 1.5 Compile Model

```
[8]: tf.keras.backend.clear_session()
# model = SegmentationModel().prepare_model(OUTPUT_CHANNEL)
model = model_u(IMG_SHAPE, OUTPUT_CHANNEL)
model.compile(optimizer="adam", loss="sparse_categorical_crossentropy")

model.summary()
```

Model: "sequential"

| Layer (type)                         | Output Shape         | Param # |
|--------------------------------------|----------------------|---------|
| layer (Layer)                        | (None, 128, 128, 3)  | 0       |
| conv2d (Conv2D)                      | (None, 128, 128, 32) | 896     |
| conv2d_1 (Conv2D)                    | (None, 128, 128, 32) | 9248    |
| max_pooling2d (MaxPooling2D)         | (None, 64, 64, 32)   | 0       |
| conv2d_2 (Conv2D)                    | (None, 64, 64, 64)   | 18496   |
| conv2d_3 (Conv2D)                    | (None, 64, 64, 64)   | 36928   |
| max_pooling2d_1 (MaxPooling2D)       | (None, 32, 32, 64)   | 0       |
| conv2d_4 (Conv2D)                    | (None, 32, 32, 128)  | 73856   |
| conv2d_5 (Conv2D)                    | (None, 32, 32, 128)  | 147584  |
| max_pooling2d_2 (MaxPooling2D)       | (None, 16, 16, 128)  | 0       |
| conv2d_6 (Conv2D)                    | (None, 16, 16, 256)  | 295168  |
| conv2d_7 (Conv2D)                    | (None, 16, 16, 256)  | 590080  |
| conv2d_transpose (Conv2DTranspose)   | (None, 32, 32, 128)  | 131200  |
| conv2d_8 (Conv2D)                    | (None, 32, 32, 128)  | 147584  |
| conv2d_9 (Conv2D)                    | (None, 32, 32, 128)  | 147584  |
| conv2d_transpose_1 (Conv2DTranspose) | (None, 64, 64, 64)   | 32832   |
| conv2d_10 (Conv2D)                   | (None, 64, 64, 64)   | 36928   |
| conv2d_11 (Conv2D)                   | (None, 64, 64, 64)   | 36928   |

```

-----
conv2d_transpose_2 (Conv2DTr (None, 128, 128, 32)      8224
-----
conv2d_12 (Conv2D)          (None, 128, 128, 35)      10115
=====
Total params: 1,723,651
Trainable params: 1,723,651
Non-trainable params: 0
-----

```

## 1.6 Callbacks

```

[9]: from datetime import date

today = date.today()
d1 = today.strftime("%d/%m/%Y") # dd/mm/YY

# return dictionary with segmentation classes (key->number, value->name)
def labels():
    return SEGMENTATION_CLASSES_DICT

# util function for generating interactive image mask from components
def wandb_mask(bg_img, pred_mask, true_mask):
    return wandb.Image(bg_img, masks={
        "prediction" : {
            "mask_data" : pred_mask,
            "class_labels" : labels()
        },
        "ground truth" : {
            "mask_data" : true_mask,
            "class_labels" : labels()
        }
    })

# early stopping callback
early_stopping_callback = tf.keras.callbacks.EarlyStopping(monitor='loss',
    ↪patience=3)

# get epochs outputs
output_epoch_callback = tf.keras.callbacks.ModelCheckpoint(filepath='model.
    ↪{epoch:02d}-{val_loss:.2f}.h5')

# always save best model
model_checkpoint_callback = keras.callbacks.ModelCheckpoint(f"models/
    ↪2022-01-05_best_model.h5", save_best_only=True)

```

## 1.7 Train

```
[10]: import wandb
      from wandb.keras import WandbCallback

      !wandb login
```

wandb: W&B API key is configured (use `wandb login --relogin` to force relogin)

```
[11]: wandb.init(project='image-segmentation', reinit=True)

      EPOCHS = 25

      history = model.fit(trainloader,
                          epochs=EPOCHS,
                          validation_data=valloader,
                          callbacks=[early_stopping_callback, WandbCallback(),
                                   ↪output_epoch_callback, model_checkpoint_callback])

      model.save(f'models/')
      wandb.finish()
```

wandb: W&B API key is configured (use `wandb login --relogin` to force relogin)

<IPython.core.display.HTML object>

```
Epoch 1/25
735/735 [=====] - 2003s 3s/step - loss: 1.8858 -
val_loss: 1.5484
Epoch 2/25
735/735 [=====] - 2095s 3s/step - loss: 1.3117 -
val_loss: 1.1279
Epoch 3/25
735/735 [=====] - 1932s 3s/step - loss: 1.0578 -
val_loss: 0.9751
Epoch 4/25
735/735 [=====] - 1923s 3s/step - loss: 0.9475 -
val_loss: 0.9847
Epoch 5/25
735/735 [=====] - 1704s 2s/step - loss: 0.8785 -
val_loss: 0.9123
Epoch 6/25
735/735 [=====] - 5324s 7s/step - loss: 0.8299 -
val_loss: 0.8680
Epoch 7/25
735/735 [=====] - 3293s 4s/step - loss: 0.7928 -
val_loss: 0.8034
Epoch 8/25
```

```

735/735 [=====] - 1838s 3s/step - loss: 0.7466 -
val_loss: 0.8009
Epoch 9/25
735/735 [=====] - 1884s 3s/step - loss: 0.7298 -
val_loss: 0.7682
Epoch 10/25
735/735 [=====] - 1851s 3s/step - loss: 0.7017 -
val_loss: 0.8202
Epoch 11/25
735/735 [=====] - 1881s 3s/step - loss: 0.6793 -
val_loss: 0.7879
Epoch 12/25
735/735 [=====] - 1886s 3s/step - loss: 0.6585 -
val_loss: 0.7472
Epoch 13/25
735/735 [=====] - 1902s 3s/step - loss: 0.6311 -
val_loss: 0.7160
Epoch 14/25
735/735 [=====] - 1874s 3s/step - loss: 0.6161 -
val_loss: 0.7415
Epoch 15/25
735/735 [=====] - 1738s 2s/step - loss: 0.5949 -
val_loss: 0.7524
Epoch 16/25
735/735 [=====] - 6412s 9s/step - loss: 0.5871 -
val_loss: 0.7277
Epoch 17/25
735/735 [=====] - 5746s 8s/step - loss: 0.5685 -
val_loss: 0.7800
Epoch 18/25
735/735 [=====] - 4961s 7s/step - loss: 0.5516 -
val_loss: 0.7244
Epoch 19/25
735/735 [=====] - 9950s 14s/step - loss: 0.5448 -
val_loss: 0.7226
Epoch 20/25
735/735 [=====] - 11476s 16s/step - loss: 0.5220 -
val_loss: 0.7511
Epoch 21/25
593/735 [=====>...] - ETA: 20:30 - loss: 0.5193

wandb: Network error (ConnectionError), entering retry loop.

735/735 [=====] - 7164s 10s/step - loss: 0.5160 -
val_loss: 0.6998
Epoch 22/25
 72/735 [=>...] - ETA: 23:10 - loss: 0.4829

```

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
/var/folders/kh/8twngknj0j9130x9cc9xlqdm0000gn/T/ipykernel_69564/2536146075.py
↳ in <module>
      3 EPOCHS = 25
      4
----> 5 history = model.fit(trainloader,

      6             epochs=EPOCHS,
      7             validation_data=valloader,

/opt/anaconda3/envs/compvis/lib/python3.8/site-packages/wandb/integration/keras
↳ keras.py in new_v2(*args, **kwargs)
    165         for cbk in cbks:
    166             set_wandb_attrs(cbk, val_data)
--> 167         return old_v2(*args, **kwargs)
    168
    169         training_arrays.orig_fit_loop = old_arrays

/opt/anaconda3/envs/compvis/lib/python3.8/site-packages/keras/engine/training.p
↳ in fit(self, x, y, batch_size, epochs, verbose, callbacks, validation_split,
↳ validation_data, shuffle, class_weight, sample_weight, initial_epoch,
↳ steps_per_epoch, validation_steps, validation_batch_size, validation_freq,
↳ max_queue_size, workers, use_multiprocessing)
    1182             _r=1):
    1183                 callbacks.on_train_batch_begin(step)
-> 1184                 tmp_logs = self.train_function(iterator)
    1185                 if data_handler.should_sync:
    1186                     context.async_wait()

/opt/anaconda3/envs/compvis/lib/python3.8/site-packages/tensorflow/python/eager
↳ def_function.py in __call__(self, *args, **kwargs)
    883
    884         with OptionalXlaContext(self._jit_compile):
--> 885             result = self._call(*args, **kwargs)
    886
    887             new_tracing_count = self.experimental_get_tracing_count()

/opt/anaconda3/envs/compvis/lib/python3.8/site-packages/tensorflow/python/eager
↳ def_function.py in _call(self, *args, **kwargs)
    915         # In this case we have created variables on the first call, so we
↳ run the
    916         # defunned version which is guaranteed to never create variables.
--> 917         return self._stateless_fn(*args, **kwargs) # pylint:
↳ disable=not-callable
    918         elif self._stateful_fn is not None:
    919             # Release the lock early so that multiple threads can perform the
↳ call

```

```

/opt/anaconda3/envs/compvis/lib/python3.8/site-packages/tensorflow/python/eager
↪function.py in __call__(self, *args, **kwargs)
    3037         (graph_function,
    3038          filtered_flat_args) = self._maybe_define_function(args, kwargs)
-> 3039     return graph_function._call_flat(
    3040         filtered_flat_args, captured_inputs=graph_function.
↪captured_inputs) # pylint: disable=protected-access
    3041

/opt/anaconda3/envs/compvis/lib/python3.8/site-packages/tensorflow/python/eager
↪function.py in _call_flat(self, args, captured_inputs, cancellation_manager)
    1961         and executing_eagerly):
    1962         # No tape is watching; skip to running the function.
-> 1963     return self._build_call_outputs(self._inference_function.call(
    1964         ctx, args, cancellation_manager=cancellation_manager))
    1965     forward_backward = self._select_forward_and_backward_functions(

/opt/anaconda3/envs/compvis/lib/python3.8/site-packages/tensorflow/python/eager
↪function.py in call(self, ctx, args, cancellation_manager)
    589     with _InterpolateFunctionError(self):
    590         if cancellation_manager is None:
--> 591         outputs = execute.execute(
    592             str(self.signature.name),
    593             num_outputs=self._num_outputs,

/opt/anaconda3/envs/compvis/lib/python3.8/site-packages/tensorflow/python/eager
↪execute.py in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
    57     try:
    58         ctx.ensure_initialized()
---> 59     tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name,
↪op_name,
    60                                         inputs, attrs, num_outputs)
    61 except core._NotOkStatusException as e:

KeyboardInterrupt:

```

## 1.8 Prediction

```

[13]: from pycm import *

def confusion_matrix(pred_mask, true_mask):

    # 2D -> 1D Array
    pred = pred_mask.ravel()

```

```

true = true_mask.ravel()

# Create Classification metrics and plot confusion matrix
cm = ConfusionMatrix(actual_vector=true, predict_vector=pred)
cm.plot()

def calc_accuracy(pred_mask, true_mask):
    # Count correct and wrong prediction
    true_prediction = np.count_nonzero((pred_mask == true_mask))
    false_prediction = np.count_nonzero((pred_mask != true_mask))
    accuracy_percentage = true_prediction*100/(false_prediction +
    ↪true_prediction)
    # print(f"Accuracy: {true_prediction}/{false_prediction+true_prediction} =
    ↪{accuracy_percentage:.1f}%")
    return accuracy_percentage

```

```

[14]: # Load model, if already trained
from keras.models import load_model

# model_ = load_model('model.15-0.76.h5')
model = load_model('models/2022-01-05_best_model.h5')

```

```

[28]: val_img, val_mask = next(iter(tstloader))

pred_mask = model.predict(val_img)
pred_mask = np.argmax(pred_mask, axis=-1)
pred_mask = np.expand_dims(pred_mask, axis=-1)

num_accuracy = 29
accuracy = np.mean([calc_accuracy(np.squeeze(pred_mask[i],-1), np.
    ↪squeeze(val_mask[i],-1)) for i in range(num_accuracy)])
print(f"Accuracy for {num_accuracy} test images: {accuracy:.1f} %")

num_example_to_display = 5
fig, axs = plt.subplots(nrows=num_example_to_display, ncols=3,
    ↪figsize=(12,4*num_example_to_display))
for i in range(num_example_to_display):
    confusion_matrix(np.squeeze(pred_mask[i],-1), np.squeeze(val_mask[i],-1))
    axs[i][0].imshow(val_img[i]);
    axs[i][1].imshow(np.squeeze(val_mask[i],-1), cmap='gray')
    axs[i][2].imshow(np.squeeze(pred_mask[i],-1), cmap='gray')

```

Accuracy for 29 test images: 60.8 %



